

Pipelining

- doesn't help latency of single task, helps the throughput of **entire** workload
- Multiple tasks operating simultaneously using **different resources**

Five Execution Stages

IF: Instruction Fetch

- Read instruction from memory using the address in PC and put it in **IF/ID** register
- PC address is incremented by **4** and then written back to the PC for next instruction

ID: Instruction Decode and Register Read

IF/ID supplies:	ID/EX receives:
<ul style="list-style-type: none"> • Register numbers for reading two registers • 16-bit offset to be sign-extended to 32-bit • PC + 4 	<ul style="list-style-type: none"> • Data values read from register file • 32-bit immediate value • PC + 4

EX: Execute an operation or calculate an address

ID/EX supplies:	EX/MEM receives:
<ul style="list-style-type: none"> • Data values read from register file • 32-bit immediate value • PC + 4 	<ul style="list-style-type: none"> • (PC + 4) + (Imm x 4) • ALU result • isZero? signal • Data Read 2 from RegFile

MEM: Access an operand in data memory

EX/MEM supplies:	MEM/WB receives:
(PC + 4) + (Imm x 4) ALU result isZero? signal Data Read 2 from RegFile	<ul style="list-style-type: none"> • ALU result • Memory read data

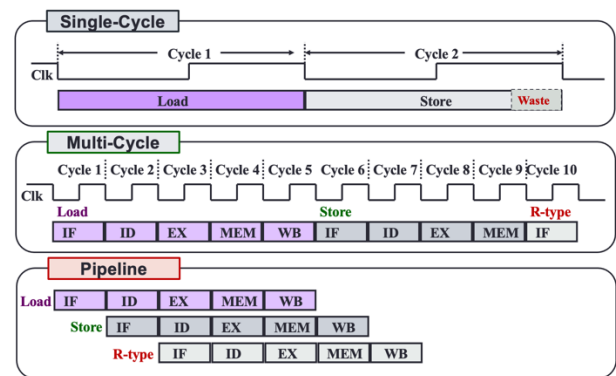
WB: Write back the result into a register

MEM/WB supplies	At the end of a cycle
ALU result Memory read data	Result is written back to register file (if applicable) using WR number*

*Pass "Write register" number from ID/EX through

EX/MEM to MEM/WB pipeline register for use in WB stage

Instructions Execution Time



Single cycle:

- $CT_{seq} = \max(\sum_1^k T_k)$
- $Time_{seq} = I \times CT_{seq}$ for I instructions

Multiple cycle:

- $CT_{multi} = \max(T_k)$
- $Time_{multi} = I \times CPI \times CT_{multi}$

CPI = num of cycles per instruction

Inst	IF	ID	EX	MEM	WB	Total
branch	X	X	X			3
R-type	X	X	X		X	4
I-type	X	X	X		X	4
lw	X	X	X	X	X	5
sw	X	X	X	X		4
j	X	X				2

* Assume that the j instruction only uses the fetch and decode stages

Pipeline cycle:

- $CT_{pipeline} = \max(T_k) + T_d$
 T_d = Overhead for pipelining, e.g. pipeline register
- Cycles needed for I instructions: $I + N - 1$
 (need $N - 1$ cycles to fill up the pipeline)
- Execution Time for I instructions:

$$Time_{pipeline} = (I + N - 1) \times CT_{pipeline}$$

Assumptions for **ideal case**:

- Every stage takes the same amount of time
- No pipeline overhead: $T_d = 0$
- $I \gg N$ (Number of instructions is much larger than number of stages)

$$Speedup_{pipeline} = \frac{Time_{seq}}{Time_{pipeline}} = \frac{I \times \sum_{k=1}^N T_k}{(I + N - 1) \times (\max(T_k) + T_d)} = \frac{I \times N \times T_1}{(I + N - 1) \times T_1}$$

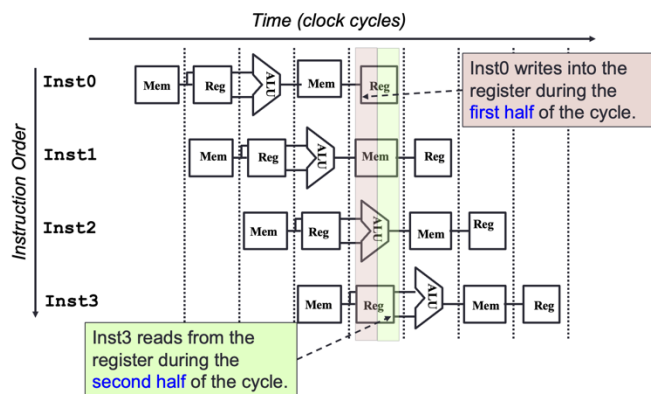
$$\approx \frac{I \times N \times T_1}{I \times T_1} = N$$

Pipeline Hazards: Problems that prevent next instruction from immediately following previous instruction

Structural Hazards

Recall that registers are very **fast memory**.

Solution: Split cycle into half; first half for writing into a register; second half for reading from a register.



Data Hazards

Register contention is the cause of **dependency**

(1) **Read-After-Write (true data dependency):** Occurs when a later instruction **reads** from the dest. register **written** by an earlier instruction

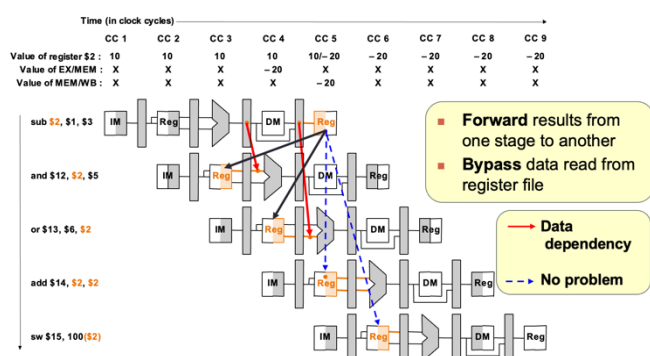
(2) **WAR: Write-after-Read** dependency

(3) **WAW: Write-after-Write** dependency

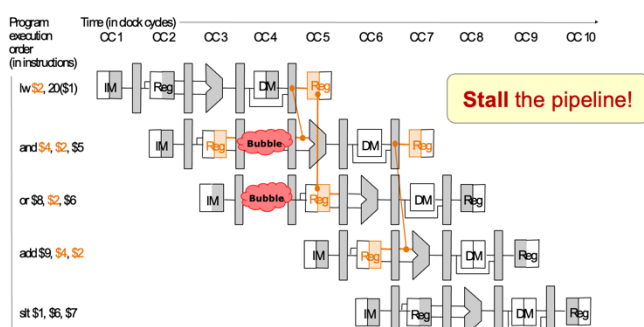
Fortunately, (2), (3) **do not** cause any pipeline hazards **unless** when instructions are executed out of program order, i.e. in Modern SuperScalar Processor

Forward

(1) No delay



(2) for **lw**, data is only ready after MEM stage may incur extra 1 cycle delay!



Control Hazards (Instruction Dependency)

Control flow is the cause of dependency


Early Branch Resolution:

- Make decision in **ID** stage instead of **MEM**
- Move branch target address calculation
- Move register comparison→not use ALU for register comparison any more (can be done with a few gates)

(1) if the register(s) involved in the comparison is

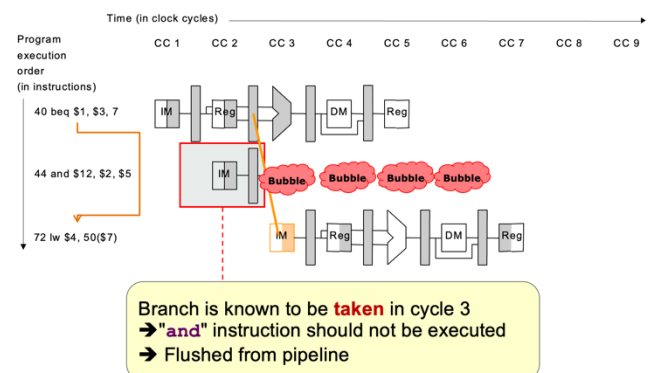
produced by **preceding** instruction: need further stall

ended up with 3 total stall cycles



Branch Prediction: Guess outcome before it is produced

- All branches are assumed to be **not taken**
- Fetch the successor instruction and start pumping it through the pipeline stages
- When the actual branch outcome is known:
Not taken→No pipeline stall
Taken: Wrong instructions in the pipeline→**Flush** successor instruction from the pipeline



MISC (From PYP)

- Predict not taken is easier to implement, as PC+4 is already computed.
- To implement predict taken, we need to compute the target branch address quickly **in the first cycle**. That requires additional hardware.

Delayed Branching:

- if the branch outcome takes X number of cycles to be known, move non-control dependent instructions into the X slots following a branch
- aka branch delay slot (the 'no-op' slots)
- instructions that will be executed regardless of the branch outcome
- In MIPS processor, Branch-Delay slot = **1** (with the early branch)