

## Basic Concepts

## PEAS Framework:

- **Performance measure:** define “goodness” of a solution
- **Environment:** define what the agent can and cannot do
- **Actuators:** outputs
- **Sensors:** inputs

## Properties of Task Environment

**Fully observable:** (vs. **Partially** observable) An agent’s sensors give it access to the **complete** state of env. at each point in time.  
**Deterministic:** (vs. **Stochastic**) Next state of env. is completely determined by current state and action executed by the agent.  
**Strategic:** env. is **deterministic** except for actions of other agents.

**Episodic:** (vs. **Sequential**) Agent’s experience is divided into atomic “episodes”. Each episode consists of agent perceiving and then performing a single action. The choice of action in each episode depends **only** on the episode itself.

**Static:** (vs. **Dynamic**) The env. is unchanged while an agent is deliberating.

**Semi-dynamic:** env. itself does not change with passage of time, but the agent’s performance score does.

**Discrete:** (vs. **Continuous**) A limited number of distinct, clearly defined percepts and actions.

**Single agent:** (vs. Multi.) An agent operating by itself in an env.

## Agents

- The agent function maps from percept histories to actions.
- An agent is completely specified by the agent function.
- A rational agent will choose actions that maximize ‘P’
- **Exploration:** Learn more about the world
- **Exploitation:** Maximize gain based on current knowledge

## Uninformed Search

**Trick:** Assume we know correct max depth for IDS/DLS

- Terminate (have sol): BFS/DLS/IDS
- Terminate (No sol): DLS
- Find Answer: BFS/DLS/IDS

**Graph Search:** Additional  $O(b^m)$  memory to store all visited nodes to avoid revisiting states.

**BFS:** Queue

```
def BFS:
    insert initial state to frontier
    while frontier is not empty:
        state = frontier.pop()
        if state is goal: return solution
        for action in actions(state):
            next state = transition(state, action)
            frontier.add(next state)
    return failure
```

- $T(n) = 1 + b + b^2 + \dots + b^d = O(b^d)$  (# nodes generated)
- $S(n) = O(b^d)$
- **Worst case:** Expand the last child in a branch
- **Complete:** Yes, if B is finite. If a solution **exists**, then the depth of the shallowest node s must be finite, so BFS must eventually search this depth. Hence, it’s complete.
  - Apply goal-test when **PUSHING** a successor state to the frontier to preserve completeness.
- **Optimal:** Generally not optimal because it simply does not take costs into consideration when determining which node to replace on the frontier.
  - Optimal is if all edge costs are equivalent.

## DFS: Stack

- $T(n) = O(b^m) \Rightarrow O(b^l)$  for DLS
- $S(n) = O(bm) \Rightarrow O(bl)$  for DLS
- **Complete:** No, when depth is infinite or can go back. There exists the possibility that DFS will faithfully yet tragically get “stuck” searching for the deepest node in an infinite-sized search tree  $\Rightarrow$  **may never find a solution.**
- **Optimal:** No, DFS simply finds the “**leftmost**” solution in the search tree without regard for path costs.

## Iterative Deepening Search (IDS)

- $T(n) = b^0 + (b^0 + b^1) + \dots + (b^0 + \dots + b^d) = (d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + 2b^{d-1} + b^d = O(b^d)$
- $S(n) = O(bd)$
- Complete: Yes
- Optimal: Yes, if step cost is the same everywhere
- Overhead =  $(\#IDS - \#DLS) / \#DLS$

## Uniform-cost Search (UCS): Priority Queue (path cost)

```
insert initial state to frontier
while frontier is not empty:
    state = frontier.pop()
    if state is goal: return solution
    for action in actions(state):
        next state = transition(state, action)
        frontier.add(next state)
return failure
```

- $T(n) = O(b^{C^*/\epsilon})$
- $S(n) = O(b^{C^*/\epsilon})$
- **Complete:** Yes, if  $\epsilon > 0$  and  $C^*$  finite. If a goal state **exists**, it must have some finite length shortest path;  $\Rightarrow$  must find shortest path.
- **Optimal:** Yes, if  $\epsilon > 0$ 
  - $C^*$  cost of optimal solution
  - $\epsilon$  minimum edge cost.  $\epsilon = 0$  may cause zero cost cycle

**Bidirectional Search:** Forward (from start) and Backward (from goal):  $\Rightarrow 2 \times O(b^{d/2}) < O(b^d)$

**Greedy Best-first Search:** Priority Queue ( $f(n) = h(n)$ )

- $T(n), S(n) = O(b^m)$ , good heuristic gives improvement
- Complete and Optimal: **No**. Particularly when using a bad heuristic function. It acts unpredictably from scenario to scenario, and can range from going straight to a goal state to acting like a badly-guided DFS and exploring all wrong areas.

**A\* Search:** Priority Queue ( $f(n) = g(n) + h(n)$ )

- $T(n), S(n) = O(b^m)$ , good heuristic gives improvement
- Complete: Yes
- Optimal: Yes
  - If  $h(n)$  **admissible** and using **tree search**.
  - if  $h(n)$  is **consistent** and using **graph search**.
  - Work with negative edge weights.
  - A\* **DO NOT** work with negative heuristics even admissible

**Admissible:** A heuristic  $h(n)$  is admissible if  $\forall n, h(n) \leq h^*(n)$

- $h^*(n)$  is the true cost to reach the goal state from  $n$ .
- **Conservative:** An admissible heuristic **never** overestimates the cost to reach the goal

**Manhattan Distance:**  $MD(x_1, y_1, x_2, y_2) = |x_1 - x_2| + |y_1 - y_2|$

Let  $p_i$  be the current location

- $\frac{1}{k} \sum_{i=1}^k MD(p_i, g) \leq \max\{MD(p_i, g)\} \leq h^*(n)$
- $\frac{1}{2} (\max\{MD(p_i, g)\} + \min\{MD(p_i, g)\}) \leq \max\{MD\} \leq h^*(n)$
- $\max\{\min\{MD(p_i, g)\}\} \leq h^*(n)$

**Consistent:**  $\forall n$ , every successor  $n'$  of  $n$  generated by any action  $a$ ,  $h(n) \leq c(n, a, n') + h(n')$ , and  $h(G) = 0$ .  $h$  is consistent  $\Rightarrow f(n') = g(n') + h(n') = g(n) + c(n, a, n') + h(n') \geq g(n) + h(n) = f(n) \Rightarrow f(n)$  is **non-decreasing** along any path

- **Theorem:** If  $h(G) = 0$ , then  $h(n)$  is consistent  $\Rightarrow$  admissible
- Admissibility **does NOT** imply consistency.

**Dominance:**  $\forall n, h_2(n) \geq h_1(n)$ , then  $h_2$  dominates  $h_1 \Leftrightarrow h_2$  is better for search if admissible.

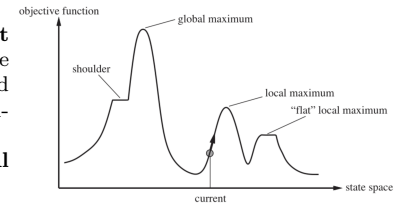
- No dominance relationship if  $h(n)$  is not admissible.

## Local Search

- **Goal:** minimizes the number of conflict violations
- Almost constant time, but **incomplete** and **sub-optimal**
- We are not interested in obtaining the solution path, but rather, reaching the goal state.
- When the search space is huge, using informed search could take a very long time.

## Hill climbing algorithm

- The algorithm **does not** maintain a search tree but only the states and the corresponding values of the objective.
- May be trapped in **local maxima**



current = initial state

**while** True:

neighbor = a highest-valued successor of current

**if** value(neighbor) <= value(current):

**return** current

current = neighbor

**Simulated Annealing:**  $P(\text{curr}, \text{next}, T) = e^{\frac{\text{value}(\text{next}) - \text{value}(\text{curr})}{T}}$

current = initial state

T = a large positive value

**while** T > 0:

**next** = a randomly selected successor of current

**if** value(**next**) > value(current): current = **next**

**else** with probability  $P(\text{current}, \text{next}, T)$ : current = **next**

decrease T

**return** current

**Theorem:** if T decreases slowly enough, simulated annealing will find a global optimum with high probability.

## Mini-Max

**def** minimax(state):

v = max\_value(state)

**return** action in successors(state) with value v

**def** max\_value(state):

**if** is\_terminal(state): **return** utility(state)

v = -inf

**for** action, next\_state in successors(state):

v = **max**(v, min\_value(next\_state))

**return** v

**def** min\_value(state):

**if** is\_terminal(state): **return** utility(state)

v = inf

**for** action, next\_state in successors(state):

v = **min**(v, max\_value(next\_state))

**return** v

- $T(n) = O(b^m)$ ,  $S(n) = O(bm)$ , with depth first exploration
- **Complete:** Yes, if tree is finite
- **Optimal:** Yes, against optimal opponent

**Alpha-beta Pruning:** Won't change the decision

```
def alpha_beta_search(state):
    v = max_value(state, -inf, inf)
    return action in successors(state) with value v
```

```
def max_value(state, alpha, beta):
    if is_terminal(state): return utility(state)
    v = -inf
    for action, next_state in successors(state):
        v = max(v, min_value(next_state, alpha, beta))
        alpha = max(alpha, v)
        if v >= beta: return v
    return v
```

```
def min_value(state, alpha, beta):
    if is_terminal(state): return utility(state)
    v = inf
    for action, next_state in successors(state):
        v = min(v, max_value(next_state, alpha, beta))
        beta = min(beta, v)
        if v <= alpha: return v
    return v
```

- Good move order improves effectiveness of pruning  $\Rightarrow O\left(b^{\frac{m}{2}}\right)$

**Evaluation Functions:** There is **no** notion of admissibility and consistency in local search and adversarial search.

### Supervised Learning

- **Regression:** predict **continuous** output
- **Classification:** predict **discrete** output

**Formalism** Assume that  $y$  is generated by  $f : x \rightarrow y$ . We want to find a hypothesis  $h: x \rightarrow \hat{y}$  (from a hypothesis class  $H$ ) s.t.  $h \approx f$  given a training set  $\{(x_1, f(x_1)), \dots, (x_N, f(x_N))\}$  We use a **learning algorithm** to find this hypothesis

**Error:** If the output of the hypothesis is a continuous value

- MSE=  $\frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$ , where  $\hat{y}_i = h(x_i)$  and  $y_i = f(x_i)$
- MAE=  $\frac{1}{N} \sum_{i=1}^N \|\hat{y}_i - y_i\|$

### Decision Tree

- **Continues-valued Attributes:** Define a discrete-valued input attribute to partition the values into a discrete set of intervals.
- **Missing Values:**
  - Assign the most common value of the attribute
  - Assign the most common value of attribute with same output
  - Assign probability to each possible value and sample
  - Drop the attribute
  - Drop the rows

```
def DTL(examples, attributes, default):
    if examples is empty: return default
    if examples have the same classification: return classification
    if attributes is empty: return mode(examples)
    best = choose_attribute(attributes, examples)
    tree = a new decision tree with root best
    for each value vi of best:
        examples.i = {rows in examples with best = vi}
        subtree = DTL(examples.i, attributes−best, mode(examples))
        add a branch to tree with label vi and subtree subtree
```

- **Accuracy:**  $\frac{TP+TN}{TP+FN+FP+TN}$
- **Recall:**  $R = TP/(TP + FN)$
- **Precision:**  $P = TP/(TP + FP)$
- **F1 Score:**

$$F1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$$

		Actual Label	
		Positive	Negative
Predicted Label	Positive	TP	FP
	Negative	FN	TN

- **Entropy:**  $I(P(v_1), \dots, P(v_n)) = - \sum_{i=1}^n P(v_i) \log_2 P(v_i)$
- **Remainder:** Entropy of children nodes  
remainder(A)= $\sum_{i=1}^v \frac{p_i+n_i}{p+n} I(\frac{p_i}{p_i+n_i}, \frac{n_i}{p_i+n_i})$
- **Information Gain:**  $IG(A) = I(\frac{p}{p+n}, \frac{n}{p+n}) - \text{remainder}(A)$

**Overfitting:** DT performance is perfect on training data, but worse on test data. DT captures data perfectly, including the noise.

**Occam's Razor:** Prefer short/simple hypotheses. In favor:

- Short/simple hypothesis that fits the data is unlikely to be coincidence
- Long/complex hypothesis that fits the data may be coincidence

Against:

- Many ways to define small sets of hypotheses (e.g., trees with prime number of nodes that uses attribute beginning with “Z”)
- Different hypotheses representations may be used instead

#### Appendix

##### Magic Entropy Number

- $I(\frac{1}{2}, \frac{1}{2}) = 1$
- $I(\frac{1}{4}, \frac{3}{4}) = 0.811$
- $I(\frac{1}{5}, \frac{4}{5}) = 0.722$
- $I(\frac{1}{6}, \frac{5}{6}) = 0.650$
- $I(\frac{1}{7}, \frac{6}{7}) = 0.592$
- $I(\frac{3}{7}, \frac{4}{7}) = 0.985$
- $I(\frac{3}{10}, \frac{7}{10}) = 0.881$
- $I(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}) = 1.585$
- $I(\frac{1}{6}, \frac{1}{3}, \frac{1}{2}) = 1.459$
- Math:  $\log_2 N = \lg N / \lg 2$
- $I(1, 0) = 0$
- $I(\frac{1}{3}, \frac{2}{3}) = 0.918$
- $I(\frac{2}{5}, \frac{3}{5}) = 0.971$
- $I(\frac{2}{7}, \frac{5}{7}) = 0.863$
- $I(\frac{3}{8}, \frac{5}{8}) = 0.954$
- $I(\frac{2}{12}, \frac{5}{12}, \frac{5}{12}) = 1.483$
- $I(\frac{1}{7}, \frac{2}{7}, \frac{4}{7}) = 1.379$

#### Decision Tree Data Table Work Space

#							
0							
1							
2							
3							
4							
5							
6							
7							
8							
9							
10							
11							
12							
13							
14							