

Basic Concepts

PEAS Framework:

- **Performance measure:** define “goodness” of a solution
- **Environment:** define what the agent can and cannot do
- **Actuators:** outputs
- **Sensors:** inputs

Properties of Task Environment

**Fully observable:** (vs. **Partially observable**) An agent’s sensors give it access to the **complete** state of env. at each point in time.

**Deterministic:** (vs. **Stochastic**) Next state of env. is completely determined by current state and action executed by the agent.

**Strategic:** env. is **deterministic** except for actions of other agents.  
**Episodic:** (vs. **Sequential**) Agent’s experience is divided into atomic “episodes”. Each episode consists of agent perceiving and then performing a single action. The choice of action in each episode depends **only** on the episode itself.

**Static:** (vs. **Dynamic**) The env. is unchanged while an agent is deliberating.

**Semi-dynamic:** env. itself does not change with passage of time, but the agent’s performance score does.

**Discrete:** (vs. **Continuous**) A limited number of distinct, clearly defined percepts and actions.

**Single agent:** (vs. Multi.) An agent operating by itself in an env. **Agents**

- The agent function maps from percept histories to actions.
- An agent is completely specified by the agent function.
- A rational agent will choose actions that maximize ‘P’
- **Exploration:** Learn more about the world
- **Exploitation:** Maximize gain based on current knowledge

Uninformed Search

**Trick:** Assume we know correct max depth for IDS/DLS

- Terminate (have sol): BFS/DLS/IDS
- Terminate (No sol): DLS
- Find Answer: BFS/DLS/IDS

**Graph Search:** Additional  $O(b^m)$  memory to store all visited nodes to avoid revisiting states.

**BFS:** Queue

- $T(n) = 1 + b + b^2 + \dots + b^d = O(b^d)$  (# nodes generated)
- $S(n) = O(b^d)$
- **Worst case:** Expand the last child in a branch
- **Complete:** Yes, if B is finite. If a solution **exists**, then the depth of the shallowest node s must be finite, so BFS must eventually search this depth. Hence, it’s complete.
  - Apply goal-test when **PUSHING** a successor state to the frontier to preserve completeness.
- **Optimal:** Generally not optimal because it simply does not take costs into consideration when determining which node to replace on the frontier.
  - Optimal is if all edge costs are equivalent.

**DFS:** Stack

- $T(n) = O(b^m) \Rightarrow O(b^l)$  for DLS
- $S(n) = O(bm) \Rightarrow O(bl)$  for DLS
- **Complete:** No, when depth is infinite or can go back. There exists the possibility that DFS will faithfully yet tragically get “stuck” searching for the deepest node in an infinite-sized search tree  $\Rightarrow$  **may never find a solution.**
- **Optimal:** No, DFS simply finds the “**leftmost**” solution in the search tree without regard for path costs.

Iterative Deepening Search (IDS)

- $T(n) = b^0 + (b^0 + b^1) + \dots + (b^0 + \dots + b^d) = (d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + 2b^{d-1} + b^d = O(b^d)$
- $S(n) = O(bd)$
- Complete: Yes
- Optimal: Yes, if step cost is the same everywhere
- Overhead = (#IDS - #DLS)/#DLS

**Uniform-cost Search (UCS):** Priority Queue (path cost)

- $T(n) = O(b^{C^*/\epsilon})$
- $S(n) = O(b^{C^*/\epsilon})$
- **Complete:** Yes, if  $\epsilon > 0$  and  $C^*$  finite. If a goal state **exists**, it must have some finite length shortest path;  $\Rightarrow$  must find shortest path.
- **Optimal:** Yes, if  $\epsilon > 0$ 
  - $C^*$  cost of optimal solution
  - $\epsilon$  minimum edge cost.  $\epsilon = 0$  may cause zero cost cycle

**Bidirectional Search:** Forward (from start) and Backward (from goal):  $\Rightarrow 2 \times O(b^{d/2}) < O(b^d)$

**Greedy Best-first Search:** Priority Queue ( $f(n) = h(n)$ )

- $T(n), S(n) = O(b^m)$ , good heuristic gives improvement
- Complete and Optimal: **No**. Particularly when using a bad heuristic function. It acts unpredictably from scenario to scenario, and can range from going straight to a goal state to acting like a badly-guided DFS and exploring all wrong areas.

**A\* Search:** Priority Queue ( $f(n) = g(n) + h(n)$ )

- $T(n), S(n) = O(b^m)$ , good heuristic gives improvement
- Complete: Yes
- Optimal: Yes
  - If  $h(n)$  **admissible** and using **tree search**.
  - if  $h(n)$  is **consistent** and using **graph search**.
  - Work with negative edge weights.
  - A\* **DO NOT** work with negative heuristics even admissible

**Admissible:** A heuristic  $h(n)$  is admissible if  $\forall n, h(n) \leq h^*(n)$

- $h^*(n)$  is the true cost to reach the goal state from  $n$ .
- **Conservative:** An admissible heuristic **never** overestimates the cost to reach the goal

**Manhattan Distance:**  $MD(x_1, y_1, x_2, y_2) = |x_1 - x_2| + |y_1 - y_2|$   
Let  $p_i$  be the current location

- $\frac{1}{k} \sum_{i=1}^k MD(p_i, g) \leq \max\{MD(p_i, g)\} \leq h^*(n)$
- $\frac{1}{2}(\max\{MD(p_i, g)\} + \min\{MD(p_i, g)\}) \leq \max\{MD\} \leq h^*(n)$
- $\max\{\min\{MD(p_i, g)\}\} \leq h^*(n)$

**Consistent:**  $\forall n$ , every successor  $n'$  of  $n$  generated by any action  $a$ ,  $h(n) \leq c(n, a, n') + h(n')$ , and  $h(G) = 0$ .  $h$  is consistent  $\Rightarrow f(n') = g(n') + h(n') = g(n) + c(n, a, n') + h(n') \geq g(n) + h(n) = f(n)$   
 $\Rightarrow f(n)$  is **non-decreasing** along any path

- **Theorem:** If  $h(G) = 0$ , then  $h(n)$  is consistent  $\Rightarrow$  admissible
- Admissibility **does NOT** imply consistency.

**Dominance:**  $\forall n, h_2(n) \geq h_1(n)$ , then  $h_2$  dominates  $h_1 \Leftrightarrow h_2$  is better for search if admissible.

- No dominance relationship if  $h(n)$  is not admissible.

Local Search

- **Goal:** minimizes the number of conflict violations
- Almost constant time, but **incomplete and sub-optimal**
- We are not interested in obtaining the solution path, but rather, reaching the goal state.
- When the search space is huge, using informed search could take a very long time.

Hill climbing algorithm

- The algorithm **does not** maintain a search tree but only the states and the corresponding values of the objective.
- May be trapped in **local maxima**
- **Simulated Annealing:**  
$$P = e^{\frac{\text{value(next)} - \text{value(curr)}}{T}}$$

**Theorem:** if T decreases slowly enough, simulated annealing will find a global optimum with high probability.

Mini-Max

- $T(n) = O(b^m), S(n) = O(bm)$ , with depth first exploration
- **Complete:** Yes, if tree is finite
- **Optimal:** Yes, against optimal opponent

**Alpha-beta Pruning:** Won’t change the decision

- Good move order improves effectiveness of pruning  $\Rightarrow O(b^{\frac{m}{2}})$

**Evaluation Functions:** There is **no** notion of admissibility and consistency in local search and adversarial search.

Supervised Learning

- **Regression:** predict **continuous** output
- **Classification:** predict **discrete** output

**Formalism** Assume that  $y$  is generated by  $f : x \rightarrow y$ . We want to find a hypothesis  $h: x \rightarrow \hat{y}$  (from a hypothesis class  $H$ ) s.t.  $h \approx f$  given a training set  $\{(x_1, f(x_1)), \dots, (x_N, f(x_N))\}$  We use a **learning algorithm** to find this hypothesis

**Error:** If the output of the hypothesis is a continuous value

- $MSE = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$ , where  $\hat{y}_i = h(x_i)$  and  $y_i = f(x_i)$
- $MAE = \frac{1}{N} \sum_{i=1}^N \|\hat{y}_i - y_i\|$

Decision Tree

- **Continues-valued Attributes:** Define a discrete-valued input attribute to partition the values into a discrete set of intervals.
- **Missing Values:**
  - Assign the most common value of the attribute
  - Assign the most common value of attribute with same output
  - Assign probability to each possible value and sample
  - Drop the attribute
  - Drop the rows

- **Accuracy:**  $\frac{TP+TN}{TP+FN+FP+TN}$
- **Precision:**  $P = TP/(TP + FP)$
- **Recall:**  $R = TP/(TP + FN)$

• **F1 Score:**  $F1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$

- **Maximize Precision/Recall**  
 $\Leftrightarrow$  **Minimize FP/FN**  
(Depends on context)

- **Entropy:**  $I(P(v_1), \dots, P(v_n)) = -\sum_{i=1}^n P(v_i) \log_2 P(v_i)$
- **Remainder:** Entropy of children nodes

remainder(A) =  $\sum_{i=1}^v \frac{p_i + n_i}{p + n} I(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i})$   
 $IG(A) = I(\frac{p}{p+n}, \frac{n}{p+n}) - \text{remainder}(A)$

**Overfitting:** DT performance is perfect on training data, but worse on test data. DT captures data perfectly, including the noise.

- **Occam’s Razor:** Prefer short/simple hypotheses. In favor:
  - Short/simple hypothesis that fits the data is unlikely to be coincidence

- Long/complex hypothesis that fits the data may be coincidence

- Many ways to define small sets of hypotheses (e.g., trees with prime number of nodes that uses attribute beginning with “Z”)
- Different hypotheses representations may be used instead

Linear Regression

**MSE:**  $J_{MSE}(w) = \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2$

- MSE loss function is **convex** for linear regression.

**Gradient Descent**  $w_j \leftarrow w_j - \gamma \frac{\partial J(w_0, w_1, \dots, w_n)}{\partial w_j}$

- $a = \frac{\partial J(w_0, w_1)}{\partial w_0}, b = \frac{\partial J(w_0, w_1)}{\partial w_1}$
- $w_0 = w_0 - \gamma a, w_1 = w_1 - \gamma b$

**Batch GD:** Consider all training examples. Much likely to get stuck in local minima

**Mini-batch GD:** Consider a subset of training examples at a time; Cheaper (Faster) per iteration; **Randomness, may escape** local minima

**Stochastic GD:** Select one random data point at a time; Cheapest (**Fastest**) per iteration; **More randomness**, may escape local minima

**Normalization:**  $x_j \leftarrow \frac{x_j - \mu_j}{\sigma_j}$  deal with features of diff. scales.

- Preserves the relationships between the data points

**Normal Equation Set**  $\frac{\partial J_{MSE}(w)}{\partial w} = 0 \Rightarrow w = (X^T X)^{-1} X^T Y$

$$X = \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_n^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix} \quad Y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \quad w = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_n \end{bmatrix}$$

	Gradient Descent	Normal Equation
$\gamma$	Need to choose	No Need
Iteration(s)	Many	None
large # of features	No Problem	Slow $(X^T X)^{-1} \rightarrow O(n^3)$
Feature scaling	May be necessary	Not necessary
Constraints		$X^T X$ should be invertible

Logistic Regression

**Logistic function (Sigmoid)**  $\sigma(z) = \frac{1}{1 + e^{-z}}$

- $\sigma(z)' = \sigma(z)(1 - \sigma(z)) \in [0, 0.25]$
- $\Pr(y = 1) = p = \sigma(h_w(x)) = \frac{1}{1 + e^{-h_w(x)}}$
- $\frac{\partial \log(p)}{\partial w_i} = - \left( \frac{1}{1 + e^{-h_w(x)}} \frac{\partial}{\partial w_i} (1 + e^{-h_w(x)}) \right) = -p(-\frac{\partial h_w(x)}{\partial w_i})e^{-h_w(x)} = p(\frac{\partial h_w(x)}{\partial w_i})(\frac{1}{p} - 1) = (1 - p) \frac{\partial h_w(x)}{\partial w_i}$
- $\frac{\partial \log(1-p)}{\partial w_i} = -\frac{\partial h_w(x)}{\partial w_i} + \frac{\partial \log(p)}{\partial w_i} = -p(\frac{\partial h_w(x)}{\partial w_i})$
- When  $h_w(x) = w^T x = \sum_{i=1}^n w_i x_i$ , we have  $\frac{\partial h_w(x)}{\partial w_i} = x_i$

**Binary cross-entropy (BCE)** Convex for logistic regression

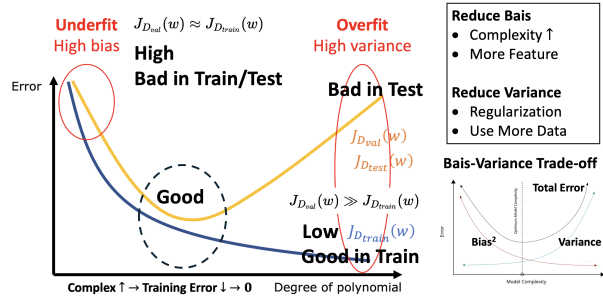
- $BCE(y, \hat{y}) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$
- $\frac{\partial BCE(y, h_w(x))}{\partial w_i} = \frac{\partial h_w(x)}{\partial w_i} (h_w(x) - y)$

• **BCE Loss**  $J_{BCE}(w) = \frac{1}{m} \sum_{i=1}^m BCE(y^{(i)}, h_w(x^{(i)}))$

- Single data with incorrect prediction  $\rightarrow \infty$  loss  $\Rightarrow$  Not suitable
- Often used for classification tasks

**Multi-class Classification:** Fit one classifier per class, fit against all other classes. Pick **highest probability**.

- Sensitivity:  $TPR = Recall = TP / (TP + FN)$
- 1-Specificity:  $FPR = FP / (FP + TN)$
- Choose threshold that have high TPR and low FPR.
- AUC: Area under ROC curve and above x-axis
- $> 0.5 \Rightarrow$  better than chance;  $\approx 1 \Rightarrow$  very accurate.
- Model is more accurate than random chance if ROC curve is above the diagonal random line.



**Evaluating the model** Accuracy and AUC-ROC metrics

**Loss Function:** Assess proximity of predict to actual

**Address Over-fitting**

- Reduce the number of features  $\Rightarrow$  High poly degree  $\rightarrow$  Low
- Regularization: Keep all features but reduce magnitude  $w_i$

**Regularization:** Parameter  $\lambda$ . For  $h_w(x) := w^T x$ , we have

• **(L2:)**  $J(w) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{i=1}^n w_j^2 \right]$

• **GD:**  $w_n \leftarrow w_n - \gamma \frac{1}{m} [\sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)}) x_n^{(i)} + \lambda w_n]$

• **Normal Equation** Works even if  $X^T X$  non-invertible if  $\lambda > 0!$

$$\lambda \sum_{j=1}^n w_j^2 \rightarrow \lambda \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \rightarrow w = \left( X^T \tilde{\lambda} \lambda + \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \right)^{-1} X^T Y$$

X:  $m \times (n+1)$   
Triangle:  $(n+1) \times (n+1)$   
Y:  $m$

• Heavily penalises larger para. Attempts to pull **all** para. small

**L1:** Use  $\lambda \sum_{i=1}^n |w_i| \Rightarrow$  Reduce # of features(set less important to 0)

**[Feature Selection]**  $\rightarrow$  Complexity  $\downarrow \rightarrow$  prevent over-fitting

**SVM**

• Decision Rule:  $w \cdot x + b \geq 0$  then +

•  $y^{(i)}(w \cdot x^{(i)} + b) - 1 = 0$  for all  $x^{(i)}$  on **margin**

$$(x^+ - x^-) \cdot \frac{w}{\|w\|} = \frac{w \cdot x^+ - w \cdot x^-}{\|w\|} = \frac{(1-b) - (-1-b)}{\|w\|} = \frac{2}{\|w\|}$$

• **Objective:**  $\max \frac{2}{\|w\|}$  s.t.  $y^{(i)}(w \cdot x^{(i)} + b) - 1 \geq 0$  [Classify correctly]  $\Leftrightarrow \min \frac{1}{2} \|w\|^2$

**Soft-Margin:** Allow misclassifications [Higher  $C \rightarrow$  less slack]

•  $\min \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi^{(i)}$  s.t.  $y^{(i)}(w \cdot x^{(i)} + b) - 1 \geq 0 - \xi^{(i)}$

$$\Rightarrow \xi^{(i)} \geq 1 - y^{(i)}(w \cdot x^{(i)} + b) \geq 0$$

•  $J(w, b) = \frac{1}{2} \|w\|^2 + C \sum \max\{0, 1 - y^{(i)}(w \cdot x^{(i)} + b)\}$

**Non-linear decision boundary:**  $w^T \phi(x) \geq 0$  then  $+$   $\rightarrow \phi$  can produce a huge number of features [Not scalable]

$\Rightarrow$  SVM: Let  $w = \sum_i \alpha^{(i)} y^{(i)} \phi(x^{(i)})$ , we have part  $\phi(x^{(i)}) \phi(x)$

• 1D:  $\phi(x) = [x, x^2]^T \Rightarrow K(u, v) = \phi(u) \phi(v) = u \cdot v$

• 2D:  $\phi(x) = [x_1, x_2, x_1 x_2, x_1^2, x_2^2, \dots]^T \Rightarrow K(u, v) = (u \cdot v)^2$

•  $n^d$  Terms:  $K(u, v) = (u \cdot v)^n$  [Kernel]

**Kernel Trick:** map training set into a different space  $\rightarrow$  linear

**Neural Network**

**Sign Function:**  $g(z) = \begin{cases} +1 & \text{if } z \geq 0 \\ -1 & \text{if } z < 0 \end{cases}$

**ReLU**  $g(x) = \max(0, x) \Rightarrow g'(x) = 1_{x>0}$

• Single-layer perceptron is a **linear** classifier

• Multi-layer perceptron can learn non-linear decision boundaries

**Perceptron Learning Algorithm**

- Initialize  $\forall_i w_i$
- Loop (until convergence or max steps reached)
  - For each  $(x^{(i)}, y^{(i)})$ , classify  $\hat{y}^{(i)} = h_w(x^{(i)}) = g(w^T x)$
  - Select one **misclassified** instance
  - Update weights:  $w \leftarrow w + \gamma(y^{(j)} - \hat{y}^{(j)})x^{(j)}$

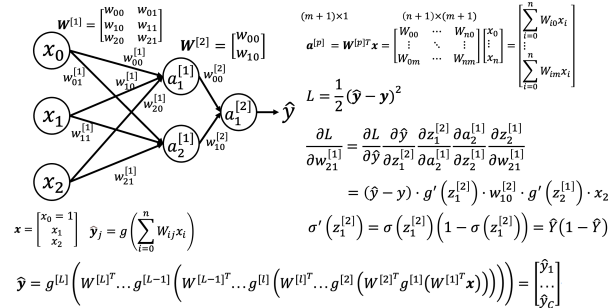
**Vanishing gradient:**  $\because \sigma'(z) \in [0, 0.25] \Rightarrow$  if we  $\times$  many small number together  $\Rightarrow w$  update small  $\Rightarrow$  **slow** convergence.

**Exploding gradient:** large gradients got multiplied again and again until it overflows. **Mitigation:**

- Using Non-saturating Activation Functions, e.g. ReLU
- Gradient Clipping

**Gradient Descent:** For signal layer  $w_i \leftarrow w_i - \gamma \frac{dL}{dw_i}$

**For Multi-Layer:**



- Without non-linear activations  $\Rightarrow$  collapses to linear ( $\hat{y} = Ax$ )
- Reorder data points help model converge faster
- Reorder may direct the model converge to a different weight (model)

**CNN**

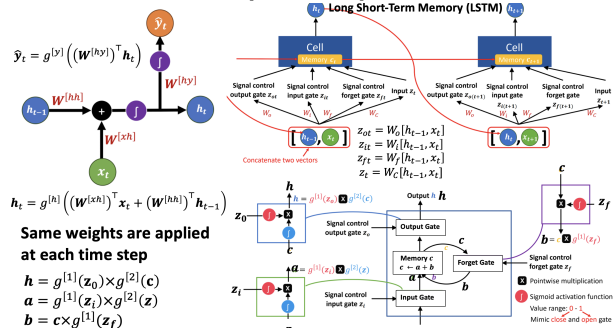
- **Convolution:** Multiply the sliding input window with **Kernel/Filter** then sum (**Not** matrix multiplication)
- **Trainable Parameter Size:**  $(W_{in} \times H_{in} \times C_{in} + 1) \times C_{out}$

**Feature Map Size:**  $\left\lfloor \frac{H-K+2P}{S} \right\rfloor + 1$

- **Max Pooling** Reduces dim, no(0) parameters to be learned
- **Fully Connected Layer** (input units  $\times$  output units) + output units
- **Batch Normalization Layer:** # of features  $\times 2$
- **Dropout** applied later to force model pay attention to all of the abstract features of the example.

**Recurrent NN:** Handling sequential data (e.g. sentence)

- One-to-many:  $T_x = 1, T_y > 1$
- Many-to-one:  $T_x > 1, T_y = 1$
- Many-to-many:  $T_x = T_y$  or  $T_x \neq T_y$



**Self-Attention**  $K^T Q = A \rightarrow A' \rightarrow H = V A'$

- **Query:** information we want to focus on.
- **Key:** info. associated with each input can be compared to query.
- **Value:** actual info. retrieved based on attention scores.

**Transformer:** Encoder-Decoder Attention

- **Query:** Generated based on previous decoder block's output
- **Key, Value:** Generated based on encoder's output
- Decoder utilize rich contextual info. provided by the encoder

**Hyperparameter Tuning**  $\Rightarrow$  Finding the best model

- Learning rate
- Epochs
- #layers  $\uparrow$
- #nodes/layer  $\uparrow$
- Activation func
- Architecture
- Kernel
- Gamma

•  $\uparrow$  hidden layers is more efficient than  $\uparrow$  nodes

**Unsupervised Learning:** Given a set of data points  $\{x^{(1)}, \dots, x^{(m)}\}$ , learn patterns in the data (No label).

- **Clustering:** identify clusters in the data
- **Dimensionality reduction:** find a lower-dimensional representation of the data

**K-Means**

- **Randomly** initialize  $K$  centroids:  $\mu_1, \dots, \mu_K$
- $\Rightarrow$  Multiple restarts of K-means give different solutions
- Repeat until convergence:
  - For  $i = 1, \dots, m$ :  $c^{(i)} \leftarrow$  index of cluster centroid ( $u_1 - u_v$ ) closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :  $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$  assigned to cluster  $k$
- No more change: converged!

It can be shown that each step in the K-Means algorithm never increases a certain loss function (the "distortion").

$$J(c^{(1)}, c^{(2)}, \dots, c^{(m)}, \mu_1, \mu_2, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

**Elbow Method:** Pick  $K$  at elbow point.  $\Rightarrow$  Heuristic method, data may not have an elbow or have multiple elbows.

**K-Medoids** Pick the **data points** that are closest to the centroids, and use them as the centroids.

**Hierarchical Clustering** Cluster Number:  $N, N-1, \dots, 1$

- Every data point is a cluster
- **Loop** (until all points are in one cluster): Find a pair of cluster that is "nearest", merge them together

**High space and time complexity:** impractical for large datasets.

**Dimensionality Reduction** SVD and PCA

**Curse of dimensionality:** # of samples to learn a hypothesis class increases **exponentially** with the # of features.

**Idea:** Change the basis of the vector to remove dependence between components (dimensions)

**Singular Value Decomposition:** For any  $n \times m$  real-valued matrix  $X$ , there exists a factorization  $X = U \Sigma V^T$  called SVD, such that

- $U$  is  $n \times m$  and has  $m$  **orthonormal** col. [**New basis**]
- $\Sigma$  is  $m \times m$  and is **diagonal** with  $\sigma_j \geq 0$  [**Basis Importance**]
- $V$  is  $m \times m$  and has  $m$  **orth.** col. and rows [**Combiner**]

**SVD Dim Reduction** Set all singular values except first  $r$  to 0. Thus  $(n \times r)(r \times r)(r \times m) = (n \times m)$  [encoder-decoder structure]

• **Data Matrix:**  $X = U \Sigma V^T \approx \tilde{U} \tilde{\Sigma} \tilde{V}^T$

• **Reduced data:**  $Z := \tilde{U}^T X$

• **Reconstructed data:**  $\hat{X} := \tilde{U} Z$

$\tilde{X} = \tilde{U} Z = \tilde{U} \tilde{U}^T X = \tilde{U} \tilde{U}^T U \Sigma V^T \approx \tilde{U} \tilde{U}^T \tilde{U} \tilde{\Sigma} \tilde{V}^T = \tilde{U} \tilde{\Sigma} \tilde{V}^T \approx U \Sigma V^T = X$

**Principal Component Analysis** By setting  $r$ , we only care about the  $r$  RVs with largest variance. Capture components that **maximize** the statistical variations of the data.

- Create the covariance matrix of the data:  $Cov(X) = \frac{1}{m} \hat{X} \hat{X}^T$
- Compute SVD on  $Cov(X)$  to obtain the  $U$  matrix (new basis)

• Reduce  $r$  components to obtain  $\tilde{U}$

Note that X's singular values^2 are related to the singular values

$$(variances) \text{ in } Cov: XX^T = U \Sigma \Sigma^T (U \Sigma \Sigma^T)^T = U \Sigma V^T V \Sigma U^T = U \Sigma \Sigma U^T = U \begin{bmatrix} \sigma_1^2 & & \\ & \sigma_r^2 & \\ & & \ddots \\ & & & \sigma_m^2 \end{bmatrix} U^T$$

**Task:** Retain at least 99% of variance in the data:

- Choose minimum  $r$ , such that  $\sum_{i=1}^r \frac{\sigma_i^2}{\sum_{i=1}^m \sigma_i^2} \geq 0.99$
- Closeness of original and reconstructed data points  $\frac{\sum_{i=1}^m \|\hat{x}^{(i)} - \tilde{x}^{(i)}\|^2}{\sum_{i=1}^m \|\hat{x}^{(i)}\|^2} \leq 0.01$

**Appendix**

**Magic Entropy Number**

- $I(\frac{1}{2}, \frac{1}{2}) = 1$
- $I(\frac{1}{4}, \frac{3}{4}) = 0.811$
- $I(\frac{1}{5}, \frac{4}{5}) = 0.722$
- $I(\frac{1}{6}, \frac{5}{6}) = 0.650$
- $I(\frac{1}{7}, \frac{6}{7}) = 0.592$
- $I(\frac{1}{8}, \frac{7}{8}) = 0.585$
- $I(\frac{3}{10}, \frac{7}{10}) = 0.881$
- $I(\frac{1}{3}, \frac{2}{3}, \frac{1}{3}) = 1.585$
- $I(\frac{1}{6}, \frac{1}{3}, \frac{1}{2}) = 1.459$
- Math:  $\log_2 N = \lg N / \lg 2$
- $I(1, 0) = 0$
- $I(\frac{1}{3}, \frac{2}{3}) = 0.918$
- $I(\frac{2}{5}, \frac{3}{5}) = 0.971$
- $I(\frac{2}{7}, \frac{5}{7}) = 0.863$
- $I(\frac{2}{8}, \frac{6}{8}) = 0.954$
- $I(\frac{2}{12}, \frac{5}{12}, \frac{5}{12}) = 1.483$
- $I(\frac{1}{7}, \frac{2}{7}, \frac{4}{7}) = 1.379$

\_\*\_\*\_\*\_\*\_\*\_\*\_- PLEASE DELETE THIS PAGE! \*\_\*\_\*\_\*\_\*\_\*\_-

### Information

Course: CS2109S Intro to AI and ML

Type: Final Cheat Sheet

Date: December 4, 2024

Author: QIU JINHANG

Link: <https://github.com/jhqiu21/Notes>

\_\*\_\*\_\*\_\*\_\*\_\*\_- PLEASE DELETE THIS PAGE! \*\_\*\_\*\_\*\_\*\_\*\_-